

```
# shield/response_evaluator.py
```

```
import re
```

```
from typing import List, Optional, Dict
```

```
from sentence_transformers import SentenceTransformer, util
```

```
import yaml
```

```
import json
```

```
import hashlib
```

```
class ResponseRuleMatch:
```

```
    id: str
```

```
    description: str
```

```
    severity: str
```

```
    score: float # Risk score for the match
```

```
class ResponseEvaluationResult:
```

```
    is_safe: bool
```

```
    risk_score: float = 0.0
```

```
    reason: Optional[str] = None
```

```
    flagged_rules: List[ResponseRuleMatch] = []
```

```
    filtered_response: Optional[str] = None
```

```
class ResponseEvaluator:
```

```
    def __init__(self, rules_path: str, embedding_model_name: Optional[str] =  
None):
```

```
        self._rules = self._load_rules(rules_path)
```

```
        self._embedding_model = SentenceTransformer(embedding_model_name)
    if embedding_model_name else None
```

```
        self._evaluation_history = {} # For loop detection
```

```
def _load_rules(self, rules_path: str) -> List[Dict]:
```

```
    try:
```

```
        with open(rules_path, 'r') as f:
```

```
            if rules_path.endswith(('yaml', 'yml')):
```

```
                return yaml.safe_load(f).get('response_rules', [])
```

```
            elif rules_path.endswith('.json'):
```

```
                return json.load(f).get('response_rules', [])
```

```
            else:
```

```
                raise ValueError("Unsupported response rules file format. Use YAML
or JSON.")
```

```
    except FileNotFoundError:
```

```
        print(f"Warning: Response rules file not found at {rules_path}. Response
evaluation will be less effective.")
```

```
        return []
```

```
    except (yaml.YAMLError, json.JSONDecodeError, ValueError) as e:
```

```
        print(f"Error loading response rules: {e}. Response evaluation will be
disabled.")
```

```
        return []
```

```
def evaluate_response(self, prompt: str, response: str, max_depth: int = 3) ->
ResponseEvaluationResult:
```

```
    initial_hash = hashlib.sha256(response.encode('utf-8')).hexdigest()
```

```
    return self._evaluate(prompt, response, 0, max_depth, {initial_hash})
```

```
def _evaluate(self, prompt: str, response: str, depth: int, max_depth: int,
seen_hashes: set) -> ResponseEvaluationResult:
```

```
    if depth >= max_depth:
```

```
        return ResponseEvaluationResult(is_safe=True, risk_score=0.0,
reason="Max response evaluation depth reached.")
```

```
    current_hash = hashlib.sha256(response.encode('utf-8')).hexdigest()
```

```
    if current_hash in seen_hashes:
```

```
        return ResponseEvaluationResult(is_safe=True, risk_score=0.0,
reason="Response loop detected.")
```

```
    seen_hashes.add(current_hash)
```

```
    result = ResponseEvaluationResult(is_safe=True, risk_score=0.0)
```

```
    for rule in self._rules:
```

```
        match_score = 0.0
```

```
        match = False
```

```
        reason = None
```

```
        match_type = rule.get("match_type", "keyword_in")
```

```
        pattern = rule.get("pattern")
```

```
        severity = rule.get("severity", "low").lower()
```

```
        weight = rule.get("weight", 1.0) # Keyword weight
```

```
        threshold = rule.get("threshold", 0.8) # Embedding similarity threshold
```

```
        if match_type == "regex" and pattern:
```

```
            if re.search(pattern, response, re.IGNORECASE):
```

```

        match = True

        match_score = self._severity_to_score(severity) * weight

        reason = f"Regex match: '{pattern}'"

    elif match_type == "keyword_in" and pattern:

        keywords = [pattern] if isinstance(pattern, str) else pattern

        for keyword in keywords:

            if keyword.lower() in response.lower():

                match = True

                match_score = max(match_score, self._severity_to_score(severity) *
weight)

                reason = f"Keyword match: '{keyword}'"

                break

    elif match_type == "embedding_similarity" and self._embedding_model
and pattern:

        response_embedding = self._embedding_model.encode(response)

        pattern_embedding = self._embedding_model.encode(pattern)

        similarity = util.cos_sim(response_embedding, pattern_embedding)[0]
[0].item()

        if similarity >= threshold:

            match = True

            match_score = max(match_score, similarity) # Use similarity as score

            reason = f"Embedding similarity >= {threshold}"

    if match:

        result.is_safe = False

        result.risk_score = max(result.risk_score, match_score)

        result.flagged_rules.append(ResponseRuleMatch(

```

```

        id=rule["id"],
        description=rule["description"],
        severity=severity,
        score=match_score
    ))

    if not result.reason:
        result.reason = reason

    for action in rule.get("actions", []):
        if isinstance(action, dict) and "filter" in action:
            filter_config = action["filter"]
            filter_type = filter_config.get("type")
            filter_pattern = filter_config.get("pattern")
            replacement = filter_config.get("replacement", "[FILTERED]")
            if filter_type == "regex_replace" and filter_pattern:
                response = re.sub(filter_pattern, replacement, response)
                result.filtered_response = response

    # Recursively evaluate the filtered response (if any)
    if result.filtered_response and result.filtered_response != response:
        recursive_result = self._evaluate(prompt, result.filtered_response, depth +
1, max_depth, seen_hashes)
        if recursive_result.risk_score > result.risk_score:
            result.risk_score = recursive_result.risk_score
            result.is_safe = recursive_result.is_safe
            result.reason = recursive_result.reason
            result.flagged_rules.extend(recursive_result.flagged_rules)

```

```
result.filtered_response = recursive_result.filtered_response
```

```
return result
```

```
def _severity_to_score(self, severity: str) -> float:
```

```
    if severity == "critical":
```

```
        return 1.0
```

```
    elif severity == "high":
```

```
        return 0.8
```

```
    elif severity == "medium":
```

```
        return 0.5
```

```
    elif severity == "low":
```

```
        return 0.2
```

```
    else:
```

```
        return 0.0
```

Updated Class/Method Structure:

The ResponseEvaluator class now includes:

\* `__init__(self, rules_path: str, embedding_model_name: Optional[str] = None):`

Takes the path to the response rules file and an optional embedding model name.

\* `_load_rules(self, rules_path: str) -> List[Dict]:` Loads response rules from YAML or JSON.

\* `evaluate_response(self, prompt: str, response: str, max_depth: int = 3) ->`

`ResponseEvaluationResult:` The main entry point for evaluating a response, initiating the recursive evaluation with loop detection.

\* \_evaluate(self, prompt: str, response: str, depth: int, max\_depth: int, seen\_hashes: set) -> ResponseEvaluationResult: A recursive helper function that performs the actual rule matching and risk scoring. It also handles loop detection.

\* \_severity\_to\_score(self, severity: str) -> float: A utility function to map rule severity to a numerical score.

The ResponseEvaluationResult class now includes a risk\_score (float) attribute. The ResponseRuleMatch class also includes a score (float) attribute representing the risk score for that specific rule match.

rules\_response.yaml Format and Example Rule:

A separate YAML file (rules\_response.yaml) will store the response rules. The format mirrors the prompt rules with the addition of weight for keyword-based matching and threshold for embedding similarity.

```
# shield/rules_response.yaml
```

```
response_rules:
```

```
- id: sensitive_info_phone
```

```
  description: Detects potential phone numbers
```

```
  severity: medium
```

```
  weight: 0.7
```

```
  pattern: "\\b\\d{3}[-\\.\\s]?\\d{3}[-\\.\\s]?\\d{4}\\b"
```

```
  match_type: regex
```

```
  actions:
```

```
    - flag:
```

```
      reason: "Potential phone number found"
```

```
    - filter:
```

```
      type: regex_replace
```

```
      pattern: "\\b\\d{3}[-\\.\\s]?\\d{3}[-\\.\\s]?\\d{4}\\b"
```

```
      replacement: "[REDACTED]"
```

- id: policy\_disallowed\_topic

description: Flags responses discussing disallowed topics

severity: high

pattern: "how to make explosives"

match\_type: keyword\_in

weight: 1.0

actions:

- flag:

reason: "Response discusses a disallowed topic"

- block\_response: true

- id: semantic\_similarity\_harmful

description: Detects responses semantically similar to harmful phrases

severity: critical

pattern: "incite violence"

match\_type: embedding\_similarity

threshold: 0.85

actions:

- flag:

reason: "Response semantically similar to harmful content"

- block\_response: true

Risk Evaluation Flow:

\* The evaluate\_response method in ResponseEvaluator is called with the prompt and the LLM response.

\* It initializes a ResponseEvaluationResult with a default risk\_score of 0.0 and is\_safe as True.



- \* The evaluator iterates through the loaded `response_rules`.
- \* For each rule, it performs the specified `match_type` (regex, keyword, embedding similarity).
- \* If a match occurs, a `match_score` is calculated based on the rule's severity (mapped to a numerical value) and optional weight (for keywords) or the similarity score itself (for embeddings).
- \* The `risk_score` in the `ResponseEvaluationResult` is updated to the maximum `match_score` found so far.
- \* A `ResponseRuleMatch` object is created, including the calculated score, and added to the `flagged_rules` list.
- \* Actions defined in the rule (e.g., flag, filter, block\_response) are executed.
- \* Based on the final `risk_score` and potentially predefined thresholds (e.g., log if > 0.3, block if > 0.7 in the calling `ShieldWrapper`), the `is_safe` status of the `ResponseEvaluationResult` can be determined.
- \* The `ResponseEvaluationResult` containing the `risk_score`, `is_safe` status, flagged rules with their scores, and any filtered response is returned.

#### Recursion Control Mechanism:

The `evaluate_response` method now implements loop detection using a recursive helper function `_evaluate` and a `seen_hashes` set:

- \* The initial call to `evaluate_response` calculates a SHA256 hash of the input response.
- \* The `_evaluate` function takes the response, the current evaluation depth, a `max_depth` limit, and a `seen_hashes` set as input.
- \* At the beginning of each recursive call, it checks if the current depth has reached `max_depth`. If so, it returns a safe result to prevent infinite loops.
- \* It calculates the SHA256 hash of the current response. If this hash is already present in the `seen_hashes` set, it indicates a loop (the response content hasn't changed), and a safe result is returned.
- \* The current response hash is added to the `seen_hashes` set before proceeding with rule evaluation.

- \* If a filter action modifies the response, the `_evaluate` function calls itself recursively with the `filtered_response`, an incremented depth, and the updated `seen_hashes` set.

- \* The results of the recursive call are incorporated into the main `ResponseEvaluationResult`.

This mechanism prevents infinite loops that could occur if response filtering leads to a response that triggers the same or other rules repeatedly without a change in content. The `max_depth` provides an additional safeguard.